

Homework 3

Question 7.1

Describe a situation or problem from your job, everyday life, current events, etc., for which exponential smoothing would be appropriate. What data would you need? Would you expect the value of alpha (the first smoothing parameter) to be closer to 0 or 1, and why?

My answer:

Problem: daily stock price for the company XXX

Data needed: the stock price over the last 2 or 5 years, the longer the better


I would expect the alpha to be closer to 1, which means there is not much randomness in the system. The stock price of the company should reflect the actual value of the company, when there is not many market uncertainties (trade war etc.). The price stays within an acceptable range of fluctuation. If we observed a fluctuation today, it probably means today's baseline is close to the observed data.

Question 7.2

Using the 20 years of daily high temperature data for Atlanta (July through October) from Question 6.2 (file temps.txt), build and use an exponential smoothing model to help make a judgment of whether the unofficial end of summer has gotten later over the 20 years.


Read in the data and plot it to see what the distribution looks like

```
temps<-read.table("temps.txt",header=TRUE)
temps_vec<-as.vector(unlist(temps[,2:21]))
plot(temps_vec)
```




Convert it to time series data and check the distribution

```
temps_ts<-ts(temps_vec,start=1996,frequency = 123)  
temps_ts  
plot(temps_ts)
```



```
plot(decompose(temps_ts))
```

Decomposition of additive time series



Apply the HoltWinters function to get the predicted value

```
temps_hw<-HoltWinters(temps_ts,alpha = NULL, beta= NULL, gamma =
NULL,seasonal = "multiplicative")
temps_hw

## Holt-Winters exponential smoothing with trend and multiplicative seasonal
component.
##
## Call:
## HoltWinters(x = temps_ts, alpha = NULL, beta = NULL, gamma = NULL,
seasonal = "multiplicative")
##
## Smoothing parameters:
##  alpha: 0.615003
##  beta : 0
##  gamma: 0.5495256
##
## Coefficients:
##              [,1]
## a    73.679517064
## b   -0.004362918
## s1    1.239022317
## s2    1.234344062
## s3    1.159509551
## s4    1.175247483
## s5    1.171344196
## s6    1.151038408
## s7    1.139383104
## s8    1.130484528
## s9    1.110487514
## s10   1.076242879
## s11   1.041044609
## s12   1.058139281
## s13   1.032496529
## s14   1.036257448
## s15   1.019348815
## s16   1.026754142
## s17   1.071170378
## s18   1.054819556
## s19   1.084397734
## s20   1.064605879
## s21   1.109827336
## s22   1.112670130
## s23   1.103970506
## s24   1.102771209
## s25   1.091264692
## s26   1.084518342
```


## s27	1.077914660
## s28	1.077696145
## s29	1.053788854
## s30	1.079454300
## s31	1.053481186
## s32	1.054023885
## s33	1.078221405
## s34	1.070145761
## s35	1.054891375
## s36	1.044587771
## s37	1.023285461
## s38	1.025836722
## s39	1.031075732
## s40	1.031419152
## s41	1.021827552
## s42	0.998177248
## s43	0.996049257
## s44	0.981570825
## s45	0.976510542
## s46	0.967977608
## s47	0.985788411
## s48	1.004748195
## s49	1.050965934
## s50	1.072515008
## s51	1.086532279
## s52	1.098357400
## s53	1.097158461
## s54	1.054827180
## s55	1.022866587
## s56	0.987259326
## s57	1.016923524
## s58	1.016604903
## s59	1.004320951
## s60	1.019102781
## s61	0.983848662
## s62	1.055888360
## s63	1.056122844
## s64	1.043478958
## s65	1.039475693
## s66	0.991019224
## s67	1.001437488
## s68	1.002221759
## s69	1.003949213
## s70	0.999566344
## s71	1.018636837
## s72	1.026490773
## s73	1.042507768
## s74	1.022500795
## s75	1.002503740
## s76	1.004560984

## s77	1.025536556
## s78	1.015357769
## s79	0.992176558
## s80	0.979377825
## s81	0.998058079
## s82	1.002553395
## s83	0.955429116
## s84	0.970970220
## s85	0.975543504
## s86	0.931515830
## s87	0.926764603
## s88	0.958565273
## s89	0.963250387
## s90	0.951644060
## s91	0.937362688
## s92	0.954257999
## s93	0.892485444
## s94	0.879537700
## s95	0.879946892
## s96	0.890633648
## s97	0.917134959
## s98	0.925991769
## s99	0.884247686
## s100	0.846648167
## s101	0.833696369
## s102	0.800001437
## s103	0.807934782
## s104	0.819343668
## s105	0.828571029
## s106	0.795608740
## s107	0.796609993
## s108	0.815503509
## s109	0.830111282
## s110	0.829086181
## s111	0.818367239
## s112	0.863958784
## s113	0.912057203
## s114	0.898308248
## s115	0.878723779
## s116	0.848971946
## s117	0.813891909
## s118	0.846821392
## s119	0.819121827
## s120	0.851036184
## s121	0.820416491
## s122	0.851581233
## s123	0.874038407

```
summary(temps_hw)
```

```
##           Length Class  Mode
## fitted      9348   mts    numeric
## x           2460   ts     numeric
## alpha        1    -none- numeric
## beta         1    -none- numeric
## gamma        1    -none- numeric
## coefficients 125    -none- numeric
## seasonal     1    -none- character
## SSE          1    -none- numeric
## call         6    -none- call
```

```
plot(temps_hw)
```



We can see from the plot that the predicted values (red line) agree pretty well with the observed values (black). The prediction begins at about year 1997, that's because 1996's data were used for prediction. The model works better with more data available. So it looks like the red line and black line align better at the later year.

Now let's take a look at the seasonal factors:

```
head(temps_hw$fitted)
```

```
##           xhat    level      trend  season
## [1,]  87.23653  82.87739 -0.004362918  1.052653
## [2,]  90.42182  82.15059 -0.004362918  1.100742
## [3,]  92.99734  81.91055 -0.004362918  1.135413
```

```
## [4,] 90.94030 81.90763 -0.004362918 1.110338
## [5,] 83.99917 81.93634 -0.004362918 1.025231
## [6,] 84.04496 81.93247 -0.004362918 1.025838

temps_hw_sf<-matrix(temps_hw$fitted[,4],nrow=123)
head(temps_hw_sf)

##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 1.052653 1.049468 1.120607 1.103336 1.118390 1.108172 1.140906
## [2,] 1.100742 1.099653 1.108025 1.098323 1.110184 1.116213 1.126827
## [3,] 1.135413 1.135420 1.139096 1.142831 1.143201 1.138495 1.129678
## [4,] 1.110338 1.110492 1.117079 1.125774 1.134539 1.126117 1.130758
## [5,] 1.025231 1.025233 1.044684 1.067291 1.084725 1.097239 1.115055
## [6,] 1.025838 1.025722 1.028169 1.042340 1.053954 1.067494 1.080203
##           [,8]      [,9]     [,10]     [,11]     [,12]     [,13]     [,14]
## [1,] 1.140574 1.125438 1.122063 1.161415 1.198102 1.198910 1.243012
## [2,] 1.154074 1.142187 1.131889 1.144549 1.134661 1.153433 1.165431
## [3,] 1.156092 1.165657 1.147982 1.149459 1.135756 1.153310 1.155197
## [4,] 1.137722 1.150639 1.146992 1.142497 1.150162 1.151169 1.157751
## [5,] 1.103877 1.120818 1.133733 1.132167 1.142714 1.139244 1.112909
## [6,] 1.094312 1.102680 1.092178 1.075766 1.088547 1.082185 1.103092
##           [,15]     [,16]     [,17]     [,18]     [,19]
## [1,] 1.243781 1.238435 1.300204 1.290647 1.254521
## [2,] 1.172935 1.190735 1.191956 1.219190 1.228826
## [3,] 1.157286 1.169773 1.189915 1.172309 1.169045
## [4,] 1.163844 1.159343 1.166605 1.167993 1.158956
## [5,] 1.132435 1.132045 1.145230 1.168161 1.170449
## [6,] 1.115071 1.118575 1.121598 1.134962 1.145475

temps_hw_smoothed<-matrix(temps_hw$fitted[,1],nrow=123)
```

I am using the predicted value (Xhat) and the observed value (data from original temps.txt) to see if summer has gotten later over the 20 years. When the difference between the predicted value and the observed value reached max, which means predicted value is much higher than the observed value, among the whole period, I assume that day is the summer end date (the model thought the temperature was still high, while the actual temperature already went down). Then I compare the date from year 1997 to year 2015 to see how the date varies.

```
fit<-temps_hw$fitted

y1997<- data.frame(fit = fit[1:123,1], obs = temps[,2])
y1997$dif <-y1997$fit - y1997$obs
head(y1997)

##      fit obs      dif
## 1 87.23653  98 -10.7634671
## 2 90.42182  97  -6.5781811
## 3 92.99734  97  -4.0026586
## 4 90.94030  90   0.9402979
```



```
## 5 83.99917 89 -5.0008320
## 6 84.04496 93 -8.9550380

a<-which.max(y1997$dif)
y1997[which.max(y1997$dif),]

##          fit obs      dif
## 100 102.7329 78 24.73294
```

The 100th day is the day where predicted temperature is much higher than the actual temperature.

```
y1998<- data.frame(fit = fit[124:246,1], obs = temps[,3])
y1998$dif <-y1998$fit - y1998$obs
b<-which.max(y1998$dif)

y1999<- data.frame(fit = fit[247:369,1], obs = temps[,3])
y1999$dif <-y1999$fit - y1999$obs
c<-which.max(y1999$dif)

y2000<- data.frame(fit = fit[370:492,1], obs = temps[,3])
y2000$dif <-y2000$fit - y2000$obs
d<-which.max(y2000$dif)

y2001<- data.frame(fit = fit[493:615,1], obs = temps[,3])
y2001$dif <-y2001$fit - y2001$obs
e<-which.max(y2001$dif)

y2002<- data.frame(fit = fit[616:738,1], obs = temps[,3])
y2002$dif <-y2002$fit - y2002$obs
f<-which.max(y2002$dif)

y2003<- data.frame(fit = fit[739:861,1], obs = temps[,3])
y2003$dif <-y2003$fit - y2003$obs
g<-which.max(y2003$dif)

y2004<- data.frame(fit = fit[862:984,1], obs = temps[,3])
y2004$dif <-y2004$fit - y2004$obs
h<-which.max(y2004$dif)

y2005<- data.frame(fit = fit[985:1107,1], obs = temps[,3])
y2005$dif <-y2005$fit - y2005$obs
i<-which.max(y2005$dif)

y2006<- data.frame(fit = fit[1108:1230,1], obs = temps[,3])
y2006$dif <-y2006$fit - y2006$obs
j<-which.max(y2006$dif)

y2007<- data.frame(fit = fit[1231:1353,1], obs = temps[,3])
y2007$dif <-y2007$fit - y2007$obs
```

```

k<-which.max(y2007$dif)

y2008<- data.frame(fit = fit[1354:1476,1], obs = temps[,3])
y2008$dif <-y2008$fit - y2008$obs
l<-which.max(y2008$dif)

y2009<- data.frame(fit = fit[1477:1599,1], obs = temps[,3])
y2009$dif <-y2009$fit - y2009$obs
m<-which.max(y2009$dif)

y2010<- data.frame(fit = fit[1600:1722,1], obs = temps[,3])
y2010$dif <-y2010$fit - y2010$obs
n<-which.max(y2010$dif)

y2011<- data.frame(fit = fit[1723:1845,1], obs = temps[,3])
y2011$dif <-y2011$fit - y2011$obs
o<-which.max(y2011$dif)

y2012<- data.frame(fit = fit[1846:1968,1], obs = temps[,3])
y2012$dif <-y2012$fit - y2012$obs
p<-which.max(y2012$dif)


y2013<- data.frame(fit = fit[1969:2091,1], obs = temps[,3])
y2013$dif <-y2013$fit - y2013$obs
q<-which.max(y2013$dif)

y2014<- data.frame(fit = fit[2092:2214,1], obs = temps[,3])
y2014$dif <-y2014$fit - y2014$obs
r<-which.max(y2014$dif)

y2015<- data.frame(fit = fit[2215:2337,1], obs = temps[,3])
y2015$dif <-y2015$fit - y2015$obs
s<-which.max(y2015$dif)

end<-c(a,b,c,d,e,f,g,h,i,j,k,l,m,n,o,p,q,r,s)
plot(end)

```



The plot shows the end days stays within a range of 100 -120 days from July 1st. So I would say the summer does not get later over the 20 years.

Question 8.1

Describe a situation or problem from your job, everyday life, current events, etc., for which a linear regression model would be appropriate. List some (up to 5) predictors that you might use.

My answer:

The height of an adult

predictors:

mother's height

father's height

gender

calories intake each week

height at age ten

Question 8.2

Using crime data from <http://www.statsci.org/data/general/uscrime.txt> (file uscrime.txt, description at <http://www.statsci.org/data/general/uscrime.html>), use regression (a useful R function is `lm` or `glm`) to predict the observed crime rate in a city with the following data: M = 14.0 So = 0 Ed = 10.0 Po1 = 12.0 Po2 = 15.5 LF = 0.640 M.F = 94.0 Pop = 150 NW = 1.1 U1 = 0.120 U2 = 3.6 Wealth = 3200 Ineq = 20.1 Prob = 0.04 Time = 39.0

Show your model (factors used and their coefficients), the software output, and the quality of fit.

```
crime<-
read.table("http://www.statsci.org/data/general/uscrime.txt",header=TRUE)
head(crime)

##      M So  Ed  Po1  Po2    LF   M.F Pop  NW   U1  U2 Wealth Ineq
## 1 15.1  1  9.1  5.8  5.6 0.510  95.0  33 30.1 0.108 4.1   3940 26.1
## 2 14.3  0 11.3 10.3  9.5 0.583 101.2  13 10.2 0.096 3.6   5570 19.4
## 3 14.2  1  8.9  4.5  4.4 0.533  96.9  18 21.9 0.094 3.3   3180 25.0
## 4 13.6  0 12.1 14.9 14.1 0.577  99.4 157  8.0 0.102 3.9   6730 16.7
## 5 14.1  0 12.1 10.9 10.1 0.591  98.5  18  3.0 0.091 2.0   5780 17.4
## 6 12.1  0 11.0 11.8 11.5 0.547  96.4  25  4.4 0.084 2.9   6890 12.6
##      Prob    Time Crime
## 1 0.084602 26.2011    791
## 2 0.029599 25.2999   1635
## 3 0.083401 24.3006    578
## 4 0.015801 29.9012   1969
## 5 0.041399 21.2998   1234
## 6 0.034201 20.9995    682

modell1<-lm(Crime~.,crime)
summary(modell1)

##
## Call:
## lm(formula = Crime ~ ., data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -395.74  -98.09   -6.69   112.99   512.67
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -5.984e+03  1.628e+03  -3.675 0.000893 ***
## M              8.783e+01  4.171e+01   2.106 0.043443 *
## So            -3.803e+00  1.488e+02  -0.026 0.979765
## Ed             1.883e+02  6.209e+01   3.033 0.004861 **
## Po1            1.928e+02  1.061e+02   1.817 0.078892 .
## Po2           -1.094e+02  1.175e+02  -0.931 0.358830
```

```
## LF          -6.638e+02  1.470e+03  -0.452  0.654654
## M.F         1.741e+01  2.035e+01   0.855  0.398995
## Pop         -7.330e-01  1.290e+00  -0.568  0.573845
## NW          4.204e+00  6.481e+00   0.649  0.521279
## U1          -5.827e+03  4.210e+03  -1.384  0.176238
## U2          1.678e+02  8.234e+01   2.038  0.050161 .
## Wealth      9.617e-02  1.037e-01   0.928  0.360754
## Ineq        7.067e+01  2.272e+01   3.111  0.003983 **
## Prob        -4.855e+03  2.272e+03  -2.137  0.040627 *
## Time        -3.479e+00  7.165e+00  -0.486  0.630708
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 209.1 on 31 degrees of freedom
## Multiple R-squared:  0.8031, Adjusted R-squared:  0.7078
## F-statistic: 8.429 on 15 and 31 DF,  p-value: 3.539e-07
```

I am using 0.05 as the threshold. Based on the p-value, I will keep M,Ed,Ineq,Prob,to fit a new model.

```
model2<-lm(Crime~M+Ed+Ineq+Prob,crime)
summary(model2)

##
## Call:
## lm(formula = Crime ~ M + Ed + Ineq + Prob, data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -532.97 -254.03  -55.72  137.80  960.21
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1339.35    1247.01  -1.074  0.28893
## M             35.97      53.39   0.674  0.50417
## Ed            148.61      71.92   2.066  0.04499 *
## Ineq           26.87      22.77   1.180  0.24458
## Prob        -7331.92    2560.27  -2.864  0.00651 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 347.5 on 42 degrees of freedom
## Multiple R-squared:  0.2629, Adjusted R-squared:  0.1927
## F-statistic: 3.745 on 4 and 42 DF,  p-value: 0.01077
```

Now only the Ed and Prob still remains significant. The adjusted R-squared dropped from 0.7078 to 0.1927. I will fit a new model to see what happens.

```
model3<-lm(Crime~Ed+Prob,crime)
summary(model3)
```

```
##
## Call:
## lm(formula = Crime ~ Ed + Prob, data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -650.98 -279.57  -14.06   198.00   957.48
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   517.30     588.48   0.879   0.3842
## Ed             63.67      50.26   1.267   0.2119
## Prob        -6049.00    2472.93  -2.446   0.0185 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 351.2 on 44 degrees of freedom
## Multiple R-squared:  0.2115, Adjusted R-squared:  0.1756
## F-statistic: 5.899 on 2 and 44 DF,  p-value: 0.005373
```

Now only Prob remains significant. The adjusted R-squared dropped from 0.1927 to 0.1756. So the first model with all the variables included turns out to be the best model so far. Next, I will try to use every combination to find the “best” model.

Create a NULL vector called model so we have something to add our layers to

```
model<-NULL
```

Create a vector of the dataframe column names used to build the formula.

```
vars <-names(crime)
```

Remove the response variable (it’s in the 16th column)

```
vars <-vars[-16]
```

The combn function will run every different combination of variables and then run the lm.

```
for(i in 1:length(vars)){
  xx = combn(vars,i)
  if(is.null(dim(xx))){
    fla = paste("Crime ~", paste(xx, collapse="+"))
  }
  model[[length(model)+1]]=lm(as.formula(fla),data=crime)
```

```

    } else {
      for(j in 1:dim(xx)[2]){
        fla = paste("Crime ~",
paste(xx[1:dim(xx)[1],j], collapse="+"))
model[[length(model)+1]]=lm(as.formula(fla),data=crime)
      }
    }
  }
}

```

See how many models were build using the loop above.

```

length(model)
## [1] 32767

```

Create a vector to extract AIC and BIC values from the model variable.

```

AICs <- NULL
BICs <- NULL
for(i in 1:length(model)){
  AICs[i] = AIC(model[[i]])
  BICs[i] = BIC(model[[i]])
}

```

See which models were chosen as best by AIC and BIC.

```

which(AICs==min(AICs))
## [1] 18494

which(BICs==min(BICs))
## [1] 5817

```

See which variables are in those models, and the corresponding adjusted R-squared.

```

summary(model[[18494]])
##
## Call:
## lm(formula = as.formula(fla), data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -444.70 -111.07   3.03  122.15  483.30
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -6426.10    1194.61  -5.379 4.04e-06 ***
## M              93.32      33.50   2.786 0.00828 **
## Ed            180.12      52.75   3.414 0.00153 **
## Po1           102.65      15.52   6.613 8.26e-08 ***

```

```
## M.F          22.34      13.60    1.642  0.10874
## U1           -6086.63   3339.27   -1.823  0.07622 .
## U2            187.35     72.48    2.585  0.01371 *
## Ineq          61.33     13.96    4.394  8.63e-05 ***
## Prob         -3796.03   1490.65   -2.547  0.01505 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 195.5 on 38 degrees of freedom
## Multiple R-squared:  0.7888, Adjusted R-squared:  0.7444
## F-statistic: 17.74 on 8 and 38 DF,  p-value: 1.159e-10
```

```
summary(model[[2496]])
```

```
##
## Call:
## lm(formula = as.formula(fla), data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -457.83 -109.01   -4.51  125.53  495.77
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -6280.671   1350.943  -4.649  4.15e-05 ***
## M              95.131     34.736    2.739  0.009431 **
## Ed            178.983     53.626    3.338  0.001935 **
## Po1           102.721     15.722    6.533  1.20e-07 ***
## M.F           21.191     14.569    1.454  0.154240
## U1           -6160.110   3395.044   -1.814  0.077724 .
## U2            189.483     73.930    2.563  0.014578 *
## Ineq          61.562     14.167    4.346  0.000104 ***
## Prob         -4066.106   1877.730   -2.165  0.036873 *
## Time          -1.431       5.920   -0.242  0.810259
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 198 on 37 degrees of freedom
## Multiple R-squared:  0.7892, Adjusted R-squared:  0.7379
## F-statistic: 15.39 on 9 and 37 DF,  p-value: 4.971e-10
```

```
summary(model[[5817]])
```

```
##
## Call:
## lm(formula = as.formula(fla), data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -470.68  -78.41  -19.68  133.12  556.23
##
```



```
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -5040.50      899.84  -5.602 1.72e-06 ***
## M           105.02       33.30   3.154 0.00305 **
## Ed          196.47       44.75   4.390 8.07e-05 ***
## Po1         115.02       13.75   8.363 2.56e-10 ***
## U2          89.37       40.91   2.185 0.03483 *
## Ineq        67.65       13.94   4.855 1.88e-05 ***
## Prob       -3801.84    1528.10  -2.488 0.01711 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 200.7 on 40 degrees of freedom
## Multiple R-squared:  0.7659, Adjusted R-squared:  0.7307
## F-statistic: 21.81 on 6 and 40 DF,  p-value: 3.418e-11

summary(model[[11564]])

##
## Call:
## lm(formula = as.formula(fla), data = crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -480.89  -89.12   -6.63   140.27   576.79
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4911.094     960.729  -5.112 8.79e-06 ***
## M           106.659      33.877   3.148 0.003144 **
## Ed          189.408      48.288   3.922 0.000345 ***
## Po1         115.704      13.993   8.269 4.16e-10 ***
## U2          88.720      41.364   2.145 0.038249 *
## Ineq        67.728      14.083   4.809 2.28e-05 ***
## Prob       -4249.756    1880.672  -2.260 0.029502 *
## Time        -2.310        5.538  -0.417 0.678810
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 202.8 on 39 degrees of freedom
## Multiple R-squared:  0.7669, Adjusted R-squared:  0.7251
## F-statistic: 18.33 on 7 and 39 DF,  p-value: 1.553e-10
```

From the output above, we can see the first two are the same model, and the last two are the same model. I will compare these two models with the model1, which has all the variables included.

```
library(data.table)

## Warning: package 'data.table' was built under R version 3.5.2
```

```

AIC(model1,model[[18494]],model[[5817]])

##           df      AIC
## model1      17 650.0291
## model[[18494]] 10 639.3151
## model[[5817]]   8 640.1661

BIC(model1,model[[18494]],model[[5817]])

##           df      BIC
## model1      17 681.4816
## model[[18494]] 10 657.8166
## model[[5817]]   8 654.9673

data.table(model1=0.7078,model18494=0.7444,model5817=0.7307)

##      model1 model18494 model5817
## 1: 0.7078      0.7444      0.7307

```

In conclusion, I would say model18494 is the best model I can found. The equation of the model is:

$$\text{crime} = -6426.10 + 93.32M + 180.12Ed + 102.65Po1 + 22.34M.F - 6086.63U1 + 187.35U2 + 61.33Ineq - 3796.03Prob$$

The corresponding AIC is 639.3151, BIC is 657.8166, and the adjusted R-squared is 0.7444.

Use the seleted model to find the crime rate in the city with data provided:

```

crimerate=-6426.10+93.32*14.0+180.12*10.0+102.65*12.0+22.34*94.0-
6086.63*0.120+187.35*3.6+61.33*20.1-3796.03*0.04
crimerate
## [1] 1038.296

```